

ML in Aid of Estimation: Part I

Lasso and ATE

Itamar Caspi

May 5, 2019 (updated: 2019-05-05)

Replicating this presentation

Use the **pacman** package to install and load packages:

```
if (!require("pacman"))  
  install.packages("pacman")  
  
pacman::p_load(tidyverse,  
               tidymodels,  
               hdm,  
               ggdag,  
               knitr,  
               xaringan,  
               RefManageR)
```

Outline

1. Causal Inference and Treatment Effects
2. Variable Selection Using the Lasso
3. Lasso In Aid of Causal Inference
4. Empirical Illustration using `hdm`

Causal Inference and Treatment Effects

The road not taken



Source: <https://mru.org/courses/mastering-econometrics/ceteris-paribus>

Notation

- Y_i is a random variable
- X_i is a vector of attributes
- \mathbf{X} is a design matrix

Treatment and potential outcomes (Rubin, 1974, 1977)

- Treatment

$$D_i = \begin{cases} 1, & \text{if unit } i \text{ received the treatment} \\ 0, & \text{otherwise.} \end{cases}$$

- Treatment and potential outcomes

Y_{i0} is the potential outcome for unit i with $D_i = 0$

Y_{i1} is the potential outcome for unit i with $D_i = 1$

- Observed outcome: Under the Stable Unit Treatment Value Assumption (SUTVA), The realization of unit i 's outcome is

$$Y_i = Y_{1i}D_i + Y_{0i}(1 - D_i)$$

Fundamental problem of causal inference (Holland, 1986): We cannot observe *both* Y_{1i} and Y_{0i} .

Treatment effect and observed outcomes

- Individual treatment effect: The difference between unit i 's potential outcomes:

$$\tau_i = Y_{1i} - Y_{0i}$$

- *Average treatment effect* (ATE)

$$\mathbb{E}[\tau_i] = \mathbb{E}[Y_{1i} - Y_{0i}] = \mathbb{E}[Y_{1i}] - \mathbb{E}[Y_{0i}]$$

- *Average treatment effect for the treatment group* (ATT)

$$\mathbb{E}[\tau_i | D_i = 1] = \mathbb{E}[Y_{1i} - Y_{0i} | D_i = 1] = \mathbb{E}[Y_{1i} | D_i = 1] - \mathbb{E}[Y_{0i} | D_i = 1]$$

NOTE: The complement of the treatment group is the *control* group.

Selection bias

A naive estimand for ATE is the difference between average outcomes based on treatment status

However, this might be misleading:

$$\mathbb{E}[Y_i|D_i = 1] - \mathbb{E}[Y_i|D_i = 0] = \underbrace{\mathbb{E}[Y_{1i}|D_i = 1] - \mathbb{E}[Y_{0i}|D_i = 1]}_{\text{ATT}} + \underbrace{\mathbb{E}[Y_{0i}|D_i = 1] - \mathbb{E}[Y_{0i}|D_i = 0]}_{\text{selection bias}}$$

Causal inference is mostly about eliminating selection-bias

EXAMPLE: Individuals who go to private universities probably have different characteristics than those who go to public universities.

Randomized control trial (RCT) solves selection bias

In an RCT, the treatments are randomly assigned. This means entails that D_i is *independent* of potential outcomes, namely

$$\{Y_{1i}, Y_{0i}\} \perp D_i$$

RCTs enables us to estimate ATE using the average difference in outcomes by treatment status:

$$\begin{aligned}\mathbb{E}[Y_i | D_i = 1] - \mathbb{E}[Y_i | D_i = 0] &= \mathbb{E}[Y_{1i} | D_i = 1] - \mathbb{E}[Y_0 | D_i = 0] \\ &= \mathbb{E}[Y_{1i} | D_i = 1] - \mathbb{E}[Y_{0i} | D_i = 1] \\ &= \mathbb{E}[Y_{1i} - Y_{0i} | D_i = 1] \\ &= \mathbb{E}[Y_{1i} - Y_{0i}] \\ &= \text{ATE}\end{aligned}$$

EXAMPLE: In theory, randomly assigning students to private and public universities would allow us to estimate the ATE going to private school have on future earnings. Clearly, RCT in this case is infeasible.

Estimands and regression

Assume for now that the treatment effect is constant across all individuals, i.e.,

$$\tau = Y_{1i} - Y_{0i}, \quad \forall i.$$

Accordingly, we can express Y_i as

$$\begin{aligned} Y_i &= Y_{1i}D_i + Y_{0i}(1 - D_i) \\ &= Y_{0i} + D_i(Y_{1i} - Y_{0i}), \\ &= Y_{0i} + \tau D_i, && \text{since } \tau = Y_{1i} - Y_{0i} \\ &= \mathbb{E}[Y_{0i}] + \tau D_i + Y_{0i} - \mathbb{E}[Y_{0i}], && \text{add and subtract } \mathbb{E}[Y_{0i}] \end{aligned}$$

Or more conveniently

$$Y_i = \alpha + \tau D_i + u_i,$$

where $\alpha = \mathbb{E}[Y_{0i}]$ and $u_i = Y_{0i} - \mathbb{E}[Y_{0i}]$ is the random component of Y_{0i} .


Unconfoundedness

In observational studies, treatments are not randomly assigned. (Think of $D_i = \{\text{private}, \text{public}\}$.)

In this case, identifying causal effects depended on the *Unconfoundedness* assumption (also known as "selection-on-observable"), which is defined as

$$\{Y_{1i}, Y_{0i}\} \perp D_i | X_i$$

In words: treatment assignment is independent of potential outcomes *conditional* on observed X_i , i.e., selection bias *disappears* when we control for X_I .



Adjusting for confounding factors

The most common approach for controlling for X_i is by adding them to the regression:

$$Y_i = \alpha + \tau D_i + X_i' \beta + u_i,$$

COMMENTS:


1. Strictly speaking, the above regression model is valid if we actually *believe* that the "true" model is $Y_i = \alpha + \tau D_i + X_i' \beta + u_i$.
2. If D_i is randomly assigned, adding X_i to the regression might increase the accuracy of ATE.
3. If D_i is assigned conditional on X_i (e.g., in observational settings), adding X_i to the regression eliminates selection bias.

An aside: Bad controls

- Bad controls are variables that are themselves outcome variables.
- This distinction becomes important when dealing with high-dimensional data

EXAMPLE: Occupation as control in a return to years of schooling regression.

Discovering that a person works as a developer in a high-tech firm changes things; knowing that the person does not have a college degree tells us immediately that he is likely to be highly capable.



Causal inference in high-dimensional setting

Consider again the standard "treatment effect regression":

$$Y_i = \alpha + \underbrace{\tau D_i}_{\text{low dimensional}} + \underbrace{\sum_{j=1}^k \beta_j X_{ij}}_{\text{high dimensional}} + \varepsilon_i, \quad \text{for } i = 1, \dots, n$$

Our object of interest is $\hat{\tau}$, the estimated *average treatment effect* (ATE).

In high-dimensional settings $k \gg n$.


Variable Selection Using the Lasso

Guarantees vs. guidance

- Most (if not all) of what we've done so far is based on *guidance*
 - Choosing the number of folds in CV
 - Size of the holdout set
 - Tuning parameter(s)
 - loss function
 - function class
- In causal inference, we need *guaranties*
 - variable selection
 - Confidence intervals and *p*-values
- To get guarantees, we typically need
 - Assumptions about a "true" model
 - Asymptotics $n \rightarrow \infty$, $k \rightarrow ?$

Resources on the theory of Lasso

- *Statistical Learning with Sparsity - The Lasso and Generalizations* (Hastie, Tibshirani, and Wainwright), **Chapter 11: Theoretical Results for the Lasso.** (PDF available online)
- *Statistics for High-Dimensional Data - Methods, Theory and Applications* (Bühlmann and van de Geer), **Chapter 7: Variable Selection with the Lasso.**
- *High Dimensional Statistics - A Non-Asymptotic Viewpoint* (Wainwright), **Chapter 7: Sparse Linear Models in High Dimensions**



Some notation to help you penetrate the Lasso literature

Suppose $\boldsymbol{\beta}$ is a $k \times 1$ vector with typical element β_i .

- The ℓ_0 -norm is defined as $\|\boldsymbol{\beta}\|_0 = \sum_{j=1}^k \mathbf{1}_{\{\beta_j \neq 0\}}$, i.e., the number of non-zero elements in $\boldsymbol{\beta}$.
- The ℓ_1 -norm is defined as $\|\boldsymbol{\beta}\|_1 = \sum_{j=1}^k |\beta_j|$.
- The ℓ_2 -norm is defined as $\|\boldsymbol{\beta}\|_2 = \left(\sum_{j=1}^k |\beta_j|^2 \right)^{\frac{1}{2}}$, i.e., Euclidean norm.
- The ℓ_∞ -norm is defined as $\|\boldsymbol{\beta}\|_\infty = \sup_j |\beta_j|$, i.e., the maximum entries' magnitude of $\boldsymbol{\beta}$.
- The support of $\boldsymbol{\beta}$, is defined as $S \equiv \text{supp}(\boldsymbol{\beta}) = \{j : \beta_j \neq 0, j = 1, \dots, k\}$, i.e., the subset of non-zero coefficients.
- The size of the support $s = |S|$ is the number of non-zero elements in $\boldsymbol{\beta}$, i.e., $s = \|\boldsymbol{\beta}\|_0$

Recap: Regularized linear regression

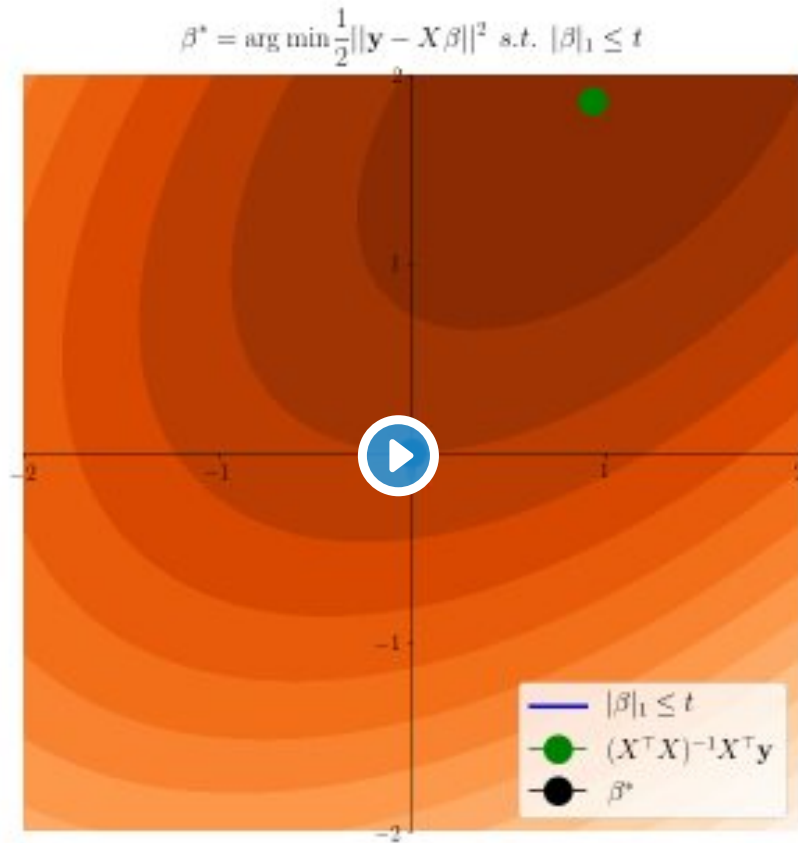
Typically, the regularized linear regression estimator is given by

$$\hat{\boldsymbol{\beta}}_{\lambda} = \arg \min_{\alpha \in \mathbb{R}, \boldsymbol{\beta} \in \mathbb{R}^k} n^{-1} \sum_{i=1}^n (Y_i - \alpha - X_i' \boldsymbol{\beta})^2 + \lambda R(\boldsymbol{\beta})$$

where

Method	$R(\boldsymbol{\beta})$
OLS	0
Subset selection	$\ \boldsymbol{\beta}\ _0$
Lasso	$\ \boldsymbol{\beta}\ _1$
Ridge	$\ \boldsymbol{\beta}\ _2^2$
Elastic Net	$\alpha \ \boldsymbol{\beta}\ _0 + (1 - \alpha) \ \boldsymbol{\beta}\ _2^2$

NOTE: We assume throughout that both Y_i and the elements in X_i have been standardized so that they have mean zero and unit variance.



Pierre Ablin
@PierreAblin



Illustration of the Lasso and its path in 2D: for t small enough, the solution is sparse!

2019 במרץ 18 - 15:50 372 ♥

Lasso: The basic setup

The linear regression model:

$$Y_i = \alpha + X_i' \boldsymbol{\beta}^0 + \varepsilon_i, \quad i = 1, \dots, n,$$

$$\mathbb{E}[\varepsilon_i X_i] = 0, \quad \alpha \in \mathbb{R}, \quad \boldsymbol{\beta}^0 \in \mathbb{R}^k.$$

Under the *exact sparsity* assumption, only a subset of variables of size $s \ll k$ is included in the model where $s \equiv \|\boldsymbol{\beta}\|_0$ is the sparsity index.

$$\underbrace{\mathbf{X}_S = (X_{(1)}, \dots, X_{(s)})}_{\text{sparse variables}}, \quad \underbrace{\mathbf{X}_{S^c} = (X_{(s+1)}, \dots, X_{(k)})}_{\text{non-sparse variables}}$$

where S is the subset of active predictors, $\mathbf{X}_S \in \mathbb{R}^{n \times s}$ corresponds to the subset of covariates that are in the sparse set, and $\mathbf{X}_{S^c} \in \mathbb{R}^{n \times k-s}$ is the subset of the "irrelevant" non-sparse variables.

Evaluation of the Lasso

Let β^0 denote the true vector of coefficients and let $\hat{\beta}$ denote the Lasso estimator.

We can assess the quality of the Lasso in several ways:

I. Prediction quality

$$\text{Loss}_{\text{pred}}(\hat{\beta}; \beta^0) = \frac{1}{N} \|\mathbf{X}\hat{\beta} - \mathbf{X}\beta^0\|_2^2$$

II. Parameter consistency

$$\text{Loss}_{\text{param}}(\hat{\beta}; \beta^*) = \|\hat{\beta} - \beta^0\|_2^2$$

III. Support recovery (sparsistency)

Lasso as a variable selection tool

- Variable selection consistency is essential for causal inference.
 - Lasso is often used as a variable selection tool.
 - Being able to select the "true" support by Lasso relies on strong assumptions about
 - the ability to distinguish between relevant and irrelevant variables.
 - the ability to identify β .
-

Critical assumption #1: Distinguishable betas

Lower eigenvalue: the min eigenvalue λ_{\min} of the sub-matrix \mathbf{X}_S is bounded away from zero.

$$\lambda_{\min} (\mathbf{X}_S' \mathbf{X}_S / N) \geq C_{\min} > 0$$

Linear dependence between the columns of \mathbf{X}_S would make it impossible to identify the true β , even if we *knew* which variables are included in \mathbf{X}_S .

NOTE: The high-dimension's lower eigenvalue condition replaces the low-dimension's rank condition (i.e., that $\mathbf{X}'\mathbf{X}$ is invertible)

Critical assumption #2: Distinguishable active predictors

Irrepresentability condition (Zou ,2006; Zhao and Yu, 2006): There must exist some $\eta \in [0, 1)$ such that


$$\max_{j \in S^c} \left\| (\mathbf{X}'_S \mathbf{X}_S)^{-1} \mathbf{X}'_S \mathbf{x}_j \right\|_1 \leq \eta$$

INTUITION: What's inside $\|\cdot\|_1$ is like regressing \mathbf{x}_j on the variables in \mathbf{X}_S .

- When $\eta = 0$, the sparse and non-sparse variables are orthogonal to each other.
- When $\eta = 1$, we can reconstruct (some elements of) \mathbf{X}_S using \mathbf{X}_{S^c} .

Thus, the irrepresentability condition roughly states that we can distinguish the sparse variables from the non-sparse ones.

Figure 2
Selected Coefficients (Nonzero Estimates) across Ten LASSO Regressions



Source: Mullainathan and Spiess (JEP 2017).

Some words on setting the optimal tuning parameter

- As we've seen thorough this course, it is also common to choose λ empirically, often by cross-validation, based on its predictive performance
- In causal analysis, inference and not prediction is the end goal. Moreover, these two objectives often contradict each other (bias vs. variance)
- Optimally, the choice of λ should provide guarantees about the performance of the model.
- Roughly speaking, when it comes to satisfying sparsistency, λ is set such that it selects non-zero β 's with high probability.

Lasso in Aid of Causal Inference

"Naive" implementation of the Lasso

Run glmnet

```
glmnet(Y ~ DX)
```

where DX is the feature matrix which includes X_i and D_i .

The estimated coefficients are:

$$\left(\hat{\alpha}, \hat{\tau}, \hat{\beta}'\right)' = \arg \min_{\alpha, \tau \in \mathbb{R}, \beta \in \mathbb{R}^{k+1}} \sum_{i=1}^n \left(Y_i - \alpha - \tau D_i - \beta' X_i\right)^2 + \lambda \left(|\tau| + \sum_{j=1}^k |\beta_j|\right)$$

PROBLEMS:

1. Both $\hat{\tau}$ and $\hat{\beta}$ are biased towards zero (shrinkage).
2. Lasso might drop D_i , i.e., shrink $\hat{\tau}$ to zero. Can also happen to relevant confounding factors.
3. How to choose λ ?

Toward a solution

OK, lets keep D_i in:

$$\left(\hat{\alpha}, \hat{\tau}, \hat{\beta}'\right)' = \arg \min_{\alpha, \tau \in \mathbb{R}, \beta \in \mathbb{R}^k} \sum_{i=1}^n \left(Y_i - \alpha - \tau D_i - \beta' X_i\right)^2 + \lambda \left(\sum_{j=1}^k |\beta_j|\right)$$

Then, *debias* the results using "Post-Lasso", i.e, use Lasso for variable selection and then run OLS with the selected variables.

PROBLEMS:

1. How to choose λ ?
2. *Omitted variable bias*: The Lasso might drop features that are correlated with D_i because they are "bad" predictor of Y_i .

Problem solved?

What can go wrong? Distribution of $\sqrt{n}(\hat{\alpha} - \alpha)$ is not what you think

$$y_i = d_i\alpha + x_i\beta + \epsilon_i, \quad d_i = x_i\gamma + v_i$$


$$\alpha = 0, \quad \beta = .2, \quad \gamma = .8,$$

$$n = 100$$

$$\epsilon_i \sim N(0, 1)$$

$$(d_i, x_i) \sim N\left(0, \begin{bmatrix} 1 & .8 \\ .8 & 1 \end{bmatrix}\right)$$

► selection done by
Lasso



Reject $H_0 : \alpha = 0$ (the truth) of no effect about 50% of the time

Source: <https://stuff.mit.edu/~vchern/papers/Chernozhukov-Saloniki.pdf>

Solution: Double-selection Lasso (Belloni, et al., REStud 2013)

First step: Regress Y_i on X_i and D_i on X_i :

$$\hat{\gamma} = \arg \min_{\gamma \in \mathbb{R}^{p+1}} \sum_{i=1}^n (Y_i - \gamma' X_i)^2 + \lambda_{\gamma} \left(\sum_{j=2}^p |\gamma_j| \right)$$
$$\hat{\delta} = \arg \min_{\delta \in \mathbb{R}^{q+1}} \sum_{i=1}^n (D_i - \delta' X_i)^2 + \lambda_{\delta} \left(\sum_{j=2}^q |\delta_j| \right)$$

Second step: Refit the model by OLS and include the \mathbf{X} 's that are significant predictors of Y_i and D_i .

Third step: Proceed to inference using standard confidence intervals.

The Tuning parameter λ is set such that the non-sparse coefficients are correctly selected with high probability.

Does it work?

Double Selection Works

$$y_i = d_i\alpha + x_i\beta + \epsilon_i, \quad d_i = x_i\gamma + v_i$$


$$\alpha = 0, \quad \beta = .2, \quad \gamma = .8,$$

$$n = 100$$

$$\epsilon_i \sim N(0, 1)$$

$$(d_i, x_i) \sim N\left(0, \begin{bmatrix} 1 & .8 \\ .8 & 1 \end{bmatrix}\right)$$

► double selection
done by Lasso



Reject $H_0 : \alpha = 0$ (the truth) about 5% of the time (nominal size = 5%)

Source: <https://stuff.mit.edu/~vchern/papers/Chernozhukov-Saloniki.pdf>

Statistical inference

Uniform Validity of the Double Selection

Theorem (Belloni, Chernozhukov, Hansen: WC 2010, ReStud 2013)

Uniformly within a class of approximately sparse models with restricted isometry conditions

$$\sigma_n^{-1} \sqrt{n}(\hat{\alpha} - \alpha_0) \rightarrow_d N(0, 1),$$

where σ_n^2 is conventional variance formula for least squares. Under homoscedasticity, semi-parametrically efficient.

- ▶ Model selection mistakes are asymptotically negligible due to double selection.
- ▶ Analogous result also holds for *endogenous* models, see Chernozhukov, Hansen, Spindler, *Annual Review of Economics*, 2015.

Source: <https://stuff.mit.edu/~vchern/papers/Chernozhukov-Saloniki.pdf>

Intuition: Partialling-out regression

consider the following two alternatives for estimating the effect of X_{1i} (a scalar) on Y_i , while adjusting for X_{2i} :

Alternative 1: Run

$$Y_i = \alpha + \beta X_{1i} + \gamma X_{2i} + \varepsilon_i$$

Alternative 2: First, run Y_i on X_{2i} and X_{1i} on X_{2i} and keep the residuals, i.e., run

$$Y_i = \gamma_0 + \gamma_1 X_{2i} + u_i^Y, \quad \text{and} \quad X_{1i} = \delta_0 + \delta_1 X_{2i} + u_i^{X_1},$$

and keep \hat{u}_i^Y and $\hat{u}_i^{X_1}$. Next, run

$$\hat{u}_i^Y = \beta^* \hat{u}_i^{X_1} + v_i.$$

According to the **Frisch-Waugh-Lovell (FWV) Theorem**,

$$\hat{\beta} = \hat{\beta}^*.$$

Notes on the guarantees of double-selection Lasso

Approximate Sparsity Consider the following regression model:

$$Y_i = f(W_i) + \varepsilon_i = X_i' \beta^0 + r_i + \varepsilon_i, \quad 1, \dots, n$$

where r_i is the approximation error.

Under *approximate sparsity*, it is assumed that $f(W_i)$ can be approximated sufficiently well (up to r_i) by $X_i' \beta^0$, while using only a small number of non-zero coefficients.

Restricted Sparse Eigenvalue Condition (RSEC) This condition puts bounds on the number of variables outside the support the Lasso can select. Relevant for the post-lasso stage.

Regularization Event The tuning parameter λ is to a value that it selects to correct model with probability of at least p , where p is set by the user. Further assumptions regarding the quantile function of the maximal value of the gradient of the objective function at β^0 , and the error term (homoskedasticity vs. heteroskedasticity). See Belloni et al. (2012) for further details.

Further extensions of double-selection

1. Chernozhukov et al. (AER 2017): Other function classes ("Double-ML"), e.g., use random forest for $Y_i \sim X_i$ and regularized logit for $D_i \sim X_i$.
2. Instrumental variables (Belloni et al., Ecta 2012, Chernozhukov et al., AER 2015)
3. Heterogeneous treatment effects (Belloni et al., Ecta 2017)
4. Panel data (Belloni, et al., JBES 2016)

Empirical Illustration using hdm

The hdm package*

"**High-Dimensional Metrics**" (hdm) by Victor Chernozhukov, Chris Hansen, and Martin Spindler is an R package for estimation and quantification of uncertainty in high-dimensional approximately sparse models.

To install the package:

```
install.packages("hdm")
```

To load the package:

```
library(hdm)
```

[*] There now also a new Stata module named **Lassopack** that includes a rich suite of programs for regularized regression in high-dimensional setting.

Illustration: Testing for growth convergence

The standard growth convergence empirical model:

$$Y_{i,T} = \alpha_0 + \alpha_1 Y_{i,0} + \sum_{j=1}^k \beta_j X_{ij} + \varepsilon_i, \quad i = 1, \dots, n,$$

where

- $Y_{i,T}$ national growth rates in GDP per capita for the periods 1965-1975 and 1975-1985.
- $Y_{i,0}$ is the log of the initial level of GDP at the beginning of the specified decade.
- X_{ij} covariates which might influence growth.

The growth convergence hypothesis implies that $\alpha_1 < 0$.

Growth data

To test the growth convergence hypothesis, we will make use of the Barro and Lee (1994) dataset

```
data("GrowthData")
```

The data contain macroeconomic information for large set of countries over several decades. In particular,

- n = 90 countries
- k = 60 country features

Not so big...

Nevertheless, the number of covariates is large relative to the sample size \Rightarrow variable selection is important!

```
library(tidyverse)
```

```
GrowthData %>% as_tibble %>% head(2)
```

```
## # A tibble: 2 x 63
##   Outcome intercept gdpsh465 bmp11 freeop freetar h65 hm65 hf65 p65
##   <dbl> <int> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 -0.0243 1 6.59 0.284 0.153 0.0439 0.007 0.013 0.001 0.290
## 2 0.100 1 6.83 0.614 0.314 0.0618 0.019 0.032 0.007 0.91
## # ... with 53 more variables: pm65 <dbl>, pf65 <dbl>, s65 <dbl>,
## # sm65 <dbl>, sf65 <dbl>, fert65 <dbl>, mort65 <dbl>, lifee065 <dbl>,
## # gpop1 <dbl>, fert1 <dbl>, mort1 <dbl>, invsh41 <dbl>, geetot1 <dbl>,
## # geerec1 <dbl>, gde1 <dbl>, govwb1 <dbl>, govsh41 <dbl>,
## # gvxdxe41 <dbl>, high65 <dbl>, highm65 <dbl>, highf65 <dbl>,
## # highc65 <dbl>, highcm65 <dbl>, highcf65 <dbl>, human65 <dbl>,
## # humanm65 <dbl>, humanf65 <dbl>, hyr65 <dbl>, hyrm65 <dbl>,
## # hyrf65 <dbl>, no65 <dbl>, nom65 <dbl>, nof65 <dbl>, pinstab1 <dbl>,
## # pop65 <int>, worker65 <dbl>, pop1565 <dbl>, pop6565 <dbl>,
## # sec65 <dbl>, secm65 <dbl>, secf65 <dbl>, secc65 <dbl>, seccm65 <dbl>,
## # seccf65 <dbl>, syr65 <dbl>, syrm65 <dbl>, syrf65 <dbl>,
## # teapri65 <dbl>, teasec65 <dbl>, ex1 <dbl>, im1 <dbl>, xr65 <dbl>,
## # tot1 <dbl>
```

Data processing

Rename the response and "treatment" variables:

```
GrowthData <- GrowthData %>%  
  rename(YT = Outcome, Y0 = gdpsh465)
```

Transform the data to vectors and matrices (to be used in the `rlassoEffect()` function)

```
YT <- GrowthData %>% select(YT) %>% pull()  
  
Y0 <- GrowthData %>% select(Y0) %>% pull()  
  
X <- GrowthData %>%  
  select(-c("Y0", "YT")) %>%  
  as.matrix()  
  
Y0_X <- GrowthData %>%  
  select(-YT) %>%  
  as.matrix()
```

Estimation of the convergence parameter α_1

Method 1: OLS

```
ols <- lm(YT ~ ., data = GrowthData)
```

Method 2: Naive (rigorous) Lasso

```
naive_Lasso <- rlasso(x = Y0_X, y = YT)
```

Does the Lasso drop Y_0 ?

```
naive_Lasso$beta[2]
```

```
## Y0
```

```
## 0
```

Unfortunately, yes...

Tidying the results

[illegible]

Results of the convergence test

```
kable(bind_rows(ols_tbl,  
               naive_Lasso_tbl,  
               part_Lasso_tbl,  
               double_Lasso_tbl),  
       format = "html")
```

method	estimate	std.error
OLS	-0.0093780	0.0298877
Naive Lasso	NA	NA
Partialling-out Lasso	-0.0498115	0.0139364
Double-selection Lasso	-0.0500059	0.0157914

Double-selection and partialling-out yield much more precise estimates and provide support the conditional convergence hypothesis


```
slides %>% end()
```

 [Source code](#)

Selected references

Ahrens, A., Hansen, C. B., & Schaffer, M. E. (2019). lassopack: Model selection and prediction with regularized regression in Stata.

Belloni, A., D. Chen, V. Chernozhukov, and C. Hansen. 2012. Sparse Models and Methods for Optimal Instruments With an Application to Eminent Domain. *Econometrica* 80(6): 2369–2429.

Belloni, A., & Chernozhukov, V. (2013). Least squares after model selection in high-dimensional sparse models. *Bernoulli*, 19(2), 521–547.

Belloni, A., Chernozhukov, V., & Hansen, C. (2013). Inference on treatment effects after selection among high-dimensional controls. *Review of Economic Studies*, 81(2), 608–650.

Belloni, A., Chernozhukov, V., & Hansen, C. (2014). High-Dimensional Methods and Inference on Structural and Treatment Effects. *Journal of Economic Perspectives*, 28(2), 29–50.

Selected references

Chernozhukov, V., Hansen, C., & Spindler, M. (2015). Post-selection and post-regularization inference in linear models with many controls and instruments. *American Economic Review*, 105(5), 486–490.

Chernozhukov, V., Hansen, C., & Spindler, M. (2016). hdm: High-Dimensional Metrics. *The R Journal*, 8(2), 185–199.

Chernozhukov, V., Chetverikov, D., Demirer, M., Duflo, E., Hansen, C., & Newey, W. (2017). Double/debiased/Neyman machine learning of treatment effects. *American Economic Review*, 107(5), 261–265.

Mullainathan, S. & Spiess, J., 2017. Machine Learning: An Applied Econometric Approach. *Journal of Economic Perspectives*, 31(2), pp.87–106.

Selected references

Van de Geer, S. A., & Bühlmann, P. (2009). On the conditions used to prove oracle results for the lasso. *Electronic Journal of Statistics*, 3, 1360–1392.

Zhao, P., & Yu, B. (2006). On Model Selection Consistency of Lasso. *Journal of Machine Learning Research*, 7, 2541–2563.